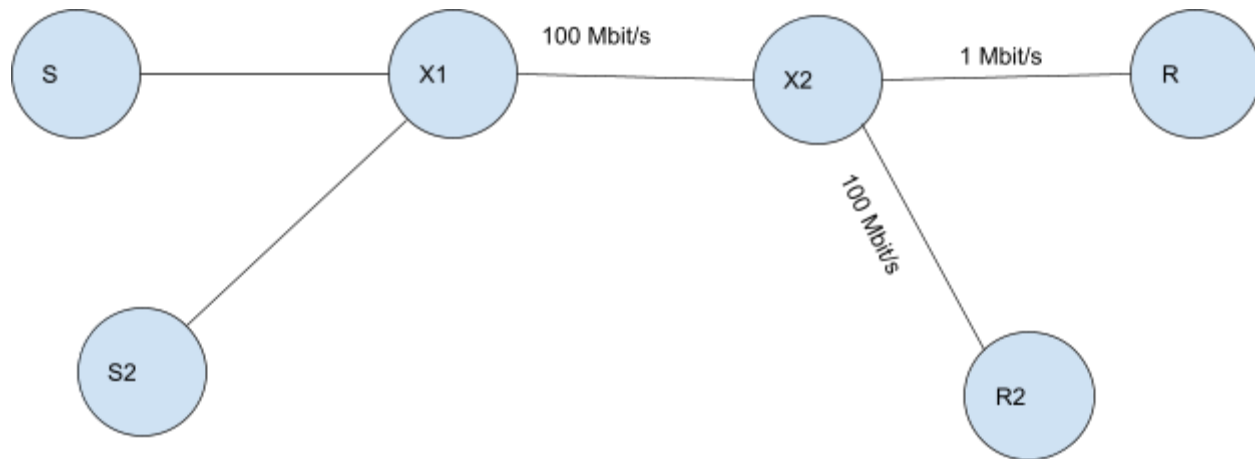


Last week: packet switching


This week: congestion control

Congestion control is resource management: assigning limited resources of link rate to flows

- TCP: flow-controlled bidirectional byte stream
  - The speed is regulated by link\_rate at the beginning. The steady state is limited by  $\min(\text{link\_rate}, \text{how fast the reader is draining the byte stream (window\_size)})$ .
- Single-flow, single-hop model  
S(sender) ----- X(router) ----- R(receiver) with  $r = 1 \text{ Gbit/s}$  and propagation delay = 1 second
  - The sender sends a datagram, still waiting for the corresponding ackno, where could the datagram be (in the sender's mind):
    - Propagating on the link
    - Waiting at the router queue (bottleneck queue)
    - Could have been received by the receiver, but ackno still on the way back
    - Or the datagram or the ackno is lost/dropped
  - Outstanding segments from the sender's perspective: [ackno, ackno + window\_size)
  - Window\_size: cap on the number of "outstanding" bytes
    - "Outstanding" means sent, not acked or judged lost
  - Q: what if the window size is really small?  
A: throughput  $\frac{1 \text{ byte}}{2 * \text{propagation delay}} = 0.5 \text{ byte/s}$  (if propagation delay = 1 s)  
(RTT = 2\*propagation delay)
  - Q: what if the window size is gigantic?  
A: maximum throughput: 1 Gbit/s and the router may run out of memory or huge queueing delay.  
We call this **congestion**.
- Congestion is bad
  - Congestion collapse (in the 1980s): receivers were advertising large window\_size and forced the router to drop lots of packet
  - Or some flows send too much, others are starved. There is an issue for fairness.
- Useful work should increase as demand increases. It's okay if the derivative is less than 1, but it should not be the case that the derivative is negative.
  - Single-flow, single-hop model would not have a collapse, since the throughput stays at 1 Gbit/s even if many packets are dropped
  - The "collapse" issue:



- 
- If there is only 1 flow:  $S2 \rightarrow R2$ , throughput would be 100 Mbit/s
- If there is two flows:  $S \rightarrow R$  and  $S2 \rightarrow R2$ , throughput would be ~51 Mbit/s if S sends at 100 Mbit/s (COLLAPSE)
- If there is two flows:  $S \rightarrow R$  and  $S2 \rightarrow R2$ , S sends at 1 Mbit/s and S2 sends at 99 Mbit/s. Throughput would be 100 Mbit/s
- The “fairness” issue:



- 
- A → a, B → b, C → c, total

15,	5,	5,	25
20,	10,	0,	30
16,	6,	4,	26
0,	0,	20,	10

max-min fair  
 max utilization/efficiency  
 proportionately fair  
 congestion collapse

- The objective: maximizing a utility function:  $\max_{\{x_r\} \in S} \sum_r U_r(x_r)$  and  $U(x) = \frac{x^{1-\alpha}}{1-\alpha}$ 
  - $\alpha = 0$ , max utilization
  - $\alpha \rightarrow 1$ , proportional fairness

- $\alpha = 2$ , min-potential-delay fairness
  - $\alpha \rightarrow \infty$ , max-min fairness
- Other objectives
  - Minimize flow completion time (of average download)
  - Minimize page load time (with many download flows)
  - Maximize “power” (= throughput / delay)
- Algorithms to prevent collapse are called “**congestion control**”
  - What is the right window size?
  - How should flows learn the window size?