

- Logistic: quiz on Wednesday

Last Time: Security

- Security Properties: Integrity, Confidentiality, Authenticity
 - Authenticity is necessary for Confidentiality

Threat Model	Mitigations / Techniques
Accidental corruption	- checksum/CRC
Adversarial modification	- Secure hash - Message Authentication Code (keyed hash)
Replay	Idempotence of messages
Eavesdropping	encryption (AE AD)
	Authenticated encryption requires a pre-established shared secret. For communication with strangers: <ul style="list-style-type: none"> - Trusted Third Party (Kerberos/Windows Active Directory): either relay the connection or Trent generates a new secret key and gives that to Alice and Bob - AKE

- peers: Alice + Bob
- eavesdropper: Eve
- adversarial modification: Mallory
- trusted third party: Trent
- AKE:
 1. Alice generates a key pair: $(public_{Alice}, private_{Alice})$
 2. Bob generates a key pair: $(public_{Bob}, private_{Bob})$
 3. Alice and Bob publishes their public keys
 4. Alice sends some x_1 to Bob, and Bob sends some x_2 to Alice
 5. Alice gets $AKE(x_1, x_2, public_{Bob}, private_{Alice}) = secret\ key$ and Bob gets $AKE(x_1, x_2, public_{Alice}, private_{Bob}) = secret\ key$, and the secret keys that Alice gets and Bob gets are the same.
 6. In addition, knowing $public_{Alice}, public_{Bob}, x_1, x_2$ does not reveal the secret key.
- But how do we know the public key of say Target?
 - Asking directly from Target does not work, since that message may be corrupted.
 - For a small number of entities, there could be a directory of public keys that were shared in a 100% secure way (e.g. an in-person meeting)
 - Or you could ask someone that you trust and you already know his/her public key
 - $sign(private_x, msg) \rightarrow signature$
 - $verify(public_x, msg, signature) \rightarrow bool$

- e.g. You are asking Keith for John's public key
 - $sign(private_{Keith}, \text{"John's public key is } \langle x \rangle \text{ according to Keith (expiring at time } t)\text{"}) \rightarrow signature_{Keith}$
 - $verify(public_{Keith}, \text{"John's public key is } \langle x \rangle \text{ according to Keith (expiring at time } t)\text{", } signature_{Keith}) \rightarrow true$. Then, this is a "certificate" that Keith verifies John's public key is $\langle x \rangle$.
 - John can store this certificate, and show this to any person that trusts Keith to prove that John is actually John.
- Firefox -----TCP----- Target @ "target.com"
 - Firefox trusts a list of certification authorities (whose public keys are programmed into Firefox)
 - When Firefox connects to "target.com", Target, to prove Target is actually Target, would provide:
 - "Hi, I'm target.com. My public key is $\langle x \rangle$. Here is a certificate from a CA you trust".
 - And a certificate: "Target come's public key is $\langle x \rangle$ according to $\langle CA \rangle$ " + $signature_{CA}$ from $private_{CA}$.
 - Firefox:
 - $verify(public_{CA}, \text{"Target come's public key is } \langle x \rangle \text{ according to } \langle CA \rangle\text{", } signature_{CA}) \rightarrow true$
 - Then Firefox and Target does *AKE* to get a shared secret key. This shared secret key is used to do *AEAD* for all following communication in the current TCP connection.
 - These steps happen as part of the TLS layer. TLS translates between plaintext and ciphertext
- Q & A
 - A: This list of CAs is common across different browsers.
 - Q: How does a CA decide to give the certificate to a specific entity?
A: CA would have an intensive verification process (back in the days), but over the time the standard has been lowered. Now it's done via domain verification: if someone can put a provided verify.txt at URL/verify.txt within 5 minutes, a CA gives the certificate. This is indeed not secure, since DNS and routing are not secure.
 - Q: What if CAs are forced to grant a certificate?
A: Certificate Transparency Log: a log of all certificates granted by CAs. Big companies monitor this